

# Machine Learning & Data Mining

## **CS/CNS/EE 155**

### Lecture 9: Conditional Random Fields

# Announcements


- Homework 5 released
  - Skeleton code available on Moodle
  - Due in 2 weeks (2/16)
- Kaggle competition closes 2/9
  - SHORT report due 2/11 via Moodle
  - Submit as a group
- Nothing due week of 2/23

# Today

- Recap of Sequence Prediction
- **Conditional Random Fields**
  - Sequential version of logistic regression
    - Analogous to how HMMs generalize Naïve Bayes
  - Discriminative sequence prediction
    - Learns to optimize  $P(y|x)$  for sequences

# Recap: Sequence Prediction


- Input:  $x = (x^1, \dots, x^M)$
- Predict:  $y = (y^1, \dots, y^M)$ 
  - Each  $y^i$  one of  $L$  labels.



- $x = \text{"Fish Sleep"}$
- $y = (N, V)$
- $x = \text{"The Dog Ate My Homework"}$
- $y = (D, N, V, D, N)$
- $x = \text{"The Fox Jumped Over The Fence"}$
- $y = (D, N, V, P, D, N)$

# Recap: General Multiclass

- $x = \text{"Fish sleep"}$
- $y = (N, V)$
- Multiclass prediction:
  - All possible length-M sequences as different class
  - $(D, D), (D, N), (D, V), (D, Adj), (D, Adv), (D, Pr)$   
 $(N, D), (N, N), (N, V), (N, Adj), (N, Adv), \dots$
- $L^M$  classes!
  - Length 2:  $6^2 = 36!$



# Recap: General Multiclass

- $x = \text{"Fish sleep"}$
- $y = (N, V)$

**POS Tags:**  
Det, Noun, Verb, Adj, Adv, Prep



$L=6$

- $M$

Can Model Everything!  
(In Theory)


SS

- $L$

Exponential Explosion in #Classes!  
(Not Tractable)

# Recap: Independent Multiclass

x="I fish often"



- Treat each word independently (assumption)
  - Independent multiclass prediction per word
  - Predict for x="I" independently
  - Predict for x="fish" independently
  - Predict for x="often" independently
  - Concatenate predictions.

Assume pronouns are nouns for simplicity.

# Recap: Independent Multiclass

$x = \text{"I fish often"}$

**POS Tags:**  
Det, Noun, Verb, Adj, Adv, Prep

$L=6$

#Classes = #POS Tags  
(6 in our example)

Solvable using standard multiclass prediction.  
But ignores context!


Assume pronouns are nouns for simplicity.



# Recap: 1<sup>st</sup> Order HMM

- $x = (x^1, x^2, x^4, x^4, \dots, x^M)$  (sequence of words)
- $y = (y^1, y^2, y^3, y^4, \dots, y^M)$  (sequence of POS tags)
- $P(x^i | y^i)$  Probability of state  $y^i$  generating  $x^i$
- $P(y^{i+1} | y^i)$  Probability of state  $y^i$  transitioning to  $y^{i+1}$
- $P(y^1 | y^0)$   $y^0$  is defined to be the Start state
- $P(\text{End} | y^M)$  Prior probability of  $y^M$  being the final state
  - Not always used

# Graphical Model Representation



Optional

$$P(x, y) = P(End \mid y^M) \prod_{i=1}^M P(y^i \mid y^{i-1}) \prod_{i=1}^M P(x^i \mid y^i)$$

# HMM Matrix Formulation

$$P(x, y) = P(END | y^M) \prod_{j=1}^M P(x^j | y^j) P(y^j | y^{j-1})$$

$$= A_{END, y^M} \prod_{j=1}^M A_{y^j y^{j-1}} O_{y^j, x^j}$$

Transition Probabilities

Emission Probabilities  
(Observation Probabilities)

# Recap: 1<sup>st</sup>-Order Sequence Models

- General multiclass:
  - Unique scoring function per entire seq.
  - Very intractable
- Independent multiclass
  - Scoring function per token, apply to each token in seq.
  - Ignores context, low accuracy
- **First-order models**
  - Scoring function per pair of tokens.
  - “Sweet spot” between fully general & ind. multiclass

# Recap: Naïve Bayes & HMMs

- Naïve Bayes:

Unique scoring function per entire seq

$$P(x, y) = P(y) \prod_{d=1}^D P(x^d | y)$$

- Hidden Markov Models:

“Naïve” Generative Independence Assumption

$$P(x, y) = P(\text{End} | y^M) \underbrace{\prod_{j=1}^M P(y^j | y^{j-1})}_{P(y)} \prod_{i=1}^M P(x^i | y^i)$$

- **HMMs  $\approx$  1<sup>st</sup> order variant of Naïve Bayes!**  
(just one interpretation...)

# Recap: Generative Models

- Joint model of  $(x, y)$ :
  - Compact & easy to train...
  - ...with ind. assumptions
    - E.g., Naïve Bayes & HMMs

$$P(x, y)$$

$\Theta$  often used to denote  
all parameters of model

- Maximize Likelihood Training:

$$\operatorname{argmax}_{\Theta} \prod_{i=1}^N P(x_i, y_i)$$

- Mismatch w/ prediction goal:
  - But hard to maximize  $P(y|x)$

$$\operatorname{argmax}_y P(y|x)$$

$$S = \{(x_i, y_i)\}_{i=1}^N$$

# Learn Conditional Prob.?

- Weird to train to maximize:

$$\operatorname{argmax}_{\Theta} \prod_{i=1}^N P(x_i, y_i)$$

$$S = \{(x_i, y_i)\}_{i=1}^N$$

- When goal should be to maximize:

$$\operatorname{argmax}_{\Theta} \prod_{i=1}^N P(y_i | x_i) = \operatorname{argmax}_{\Theta} \prod_{i=1}^N \frac{P(x_i, y_i)}{P(x_i)}$$

**Breaks independence!**

Can no longer use count statistics

~~$$P(x^d = a | y = z) = \frac{\sum_{i=1}^N 1_{[(y_i=z) \wedge (x_i^d=a)]}}{\sum_{i=1}^N 1_{[y_i=z]}}$$~~

$$p(x) = \sum_y P(x, y) = \sum_y P(y) P(x | y)$$

Both HMMs & Naïve Bayes suffer this problem!

# Learn Conditional Prob.?

- Weird to train to maximize:

$$\sum_{i=1}^N \log p(x_i | y_i)$$

In general, you should maximize the likelihood of the model you define!

So if you define joint model  $P(x,y)$ , then maximize  $P(x,y)$  on training data.

B  
C  
.

~~$$P(x^d = a | y = z) = \frac{\sum_{i=1}^N \mathbb{1}_{[y_i = z] \wedge [x_i^d = a]}}{\sum_{i=1}^N \mathbb{1}_{[y_i = z]}}$$~~

Both HMMs & Naïve Bayes suffer this problem!



# Generative vs Discriminative

- Generative Models:

Hidden Markov Models  
Naïve Bayes

- Joint Distribution:  $P(x,y)$  ← **Mismatch!**
- Uses Bayes's Rule to predict:  $\operatorname{argmax}_y P(y|x)$  ↗
- Can generate new samples  $(x,y)$

- Discriminative Models:

Conditional Random Fields  
Logistic Regression

- Conditional Distribution:  $P(y|x)$  ← **Same thing!**
- Can directly to predict:  $\operatorname{argmax}_y P(y|x)$  ↙

- Both trained via Maximum Likelihood

# First Try

(for classifying a single  $y$ )

- Model  $P(y|x)$  for every possible  $x$

$P(y=1 x)$	$x^1$	$x^2$
0.5	0	0
0.7	0	1
0.2	1	0
0.4	1	1

- Train by counting frequencies
- **Exponential in # input variables!**
  - Need to assume something... what?

# Log Linear Models!

## (Logistic Regression)

$$P(y | x) = \frac{\exp\{w_y^T x - b_y\}}{\sum_k \exp\{w_k^T x - b_k\}} \quad \begin{array}{l} x \in R^D \\ y \in \{1, 2, \dots, L\} \end{array}$$

- “Log-Linear” assumption
  - Model representation to linear in  $x$
  - Most common discriminative probabilistic model

**Prediction:**

$$\operatorname{argmax}_y P(y | x)$$

**Training:**

$$\operatorname{argmax}_{\Theta} \prod_{i=1}^N P(y_i | x_i)$$

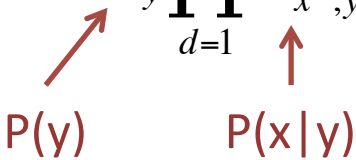
← Match! →

# Naïve Bayes vs Logistic Regression

- Naïve Bayes:

- Strong ind. assumptions
- Super easy to train...
- ...but mismatch with prediction

$$P(x, y) = A_y \prod_{d=1}^D O_{x^d, y}^d$$

  
 $P(y)$                        $P(x|y)$

- Logistic Regression:

- “Log Linear” assumption
  - Often more flexible than Naïve Bayes
- Harder to train (gradient desc.)...
- ...but matches prediction

$$P(y|x) = \frac{\exp\{w_y^T x - b_y\}}{\sum_k \exp\{w_k^T x - b_k\}}$$

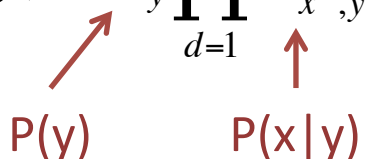
$$x \in R^D$$
$$y \in \{1, 2, \dots, L\}$$

# Naïve Bayes vs Logistic Regression

- NB has L parameters for  $P(y)$  (i.e., A)
- LR has L parameters for bias b
- NB has L\*D parameters for  $P(x|y)$  (i.e, O)
- LR has L\*D parameters for w
- **Same number of parameters!**

Naïve Bayes

$$P(x, y) = A_y \prod_{d=1}^D O_{x^d, y}^d$$



$P(y)$   $P(x|y)$

Logistic Regression

$$P(y | x) = \frac{e^{w_y^T x - b_y}}{\sum_k e^{w_k^T x - b_k}}$$

$$x \in \{0, 1\}^D$$
$$y \in \{1, 2, \dots, L\}$$

# Naïve Bayes vs Logistic Regression

## Intuition:

Both models have same “capacity”

NB spends a lot of capacity on  $P(x)$

LR spends all of capacity on  $P(y|x)$

## No Model Is Perfect!

(Especially on finite training set)

NB will trade off  $P(y|x)$  with  $P(x)$

LR will fit  $P(y|x)$  as well as possible

# Conditional Random Fields

## Sequential Version of Logistic Regression

# “Log-Linear” 1<sup>st</sup> Order Sequential Model

$$P(y | x) = \frac{1}{Z(x)} \exp \left\{ \sum_{j=1}^M \left( A_{y^j, y^{j-1}} + O_{y^j, x^j} \right) \right\}$$

$$Z(x) = \sum_{y'} \exp \{ F(y', x) \}$$

aka “Partition Function”

$$F(y, x) \equiv \sum_{j=1}^M \left( A_{y^j, y^{j-1}} + O_{y^j, x^j} \right)$$

Scoring Function

Scoring transitions

Scoring input features

$$P(y | x) = \frac{\exp \{ F(y, x) \}}{Z(x)}$$

$$\log P(y | x) = F(y, x) - \log(Z(x))$$

$y^0$  = special start state, excluding end state



- $x = \text{"Fish Sleep"}$
- $y = (N, V)$

$$P(y | x) = \frac{1}{Z(x)} \exp \left\{ \sum_{j=1}^M \left( A_{y^j, y^{j-1}} + O_{y^j, x^j} \right) \right\}$$

$A_{N,V}$  →

	$A_{N,*}$	$A_{V,*}$
$A_{*,N}$	-2	1
$A_{*,V}$	2	-2
$A_{*,Start}$	1	-1

←  $w_{V,Fish}$

	$O_{N,*}$	$O_{V,*}$
$O_{*,Fish}$	2	1
$O_{*,Sleep}$	1	0

$$P(N, V | \text{"Fish Sleep"}) = \frac{1}{Z(x)} \exp \{ A_{N,Start} + O_{N,Fish} + A_{V,N} + O_{V,Sleep} \} = \frac{1}{Z(x)} \exp \{ 4 \} \approx 0.66$$


$$Z(x) = \text{Sum} \left( \begin{array}{|c|c|} \hline y & \exp(F(y,x)) \\ \hline (N,N) & \exp(1+2-2+1) = \exp(2) \\ (N,V) & \exp(1+2+2+0) = \exp(4) \\ (V,N) & \exp(-1+1+2+1) = \exp(3) \\ (V,V) & \exp(-1+1-2+0) = \exp(-2) \\ \hline \end{array} \right)$$

- $x = \text{"Fish Sleep"}$
- $y = (N, V)$

$$P(N, V | \text{"Fish Sleep"}) = \frac{1}{Z(x)} \exp\{F(x, y)\}$$

$P(N, V | \text{"Fish Sleep"})$

\*hold other parameters fixed




$F(y, x)$

# Basic Conditional Random Field

- Directly models  $P(y|x)$ 
  - Discriminative
  - Log linear assumption
  - Same #parameters as HMM
  - 1<sup>st</sup> Order Sequential LR

CRF spends all model capacity on  $P(y|x)$ , rather than  $P(x,y)$



$$F(y, x) \equiv \sum_{j=1}^M \left( A_{y^j, y^{j-1}} + O_{y^j, x^j} \right)$$

$$P(y|x) = \frac{\exp\{F(y, x)\}}{\sum_{y'} \exp\{F(y', x)\}}$$

- **How to Predict?**
- **How to Train?**
- **Extensions?**

$$\log P(y|x) = F(y, x) - \log \left( \sum_{y'} \exp\{F(y', x)\} \right)$$


# Predict via Viterbi

$$\begin{aligned}
 \operatorname{argmax}_y P(y | x) &= \operatorname{argmax}_y \log P(y | x) = \operatorname{argmax}_y F(y, x) \\
 &= \operatorname{argmax}_y \sum_{j=1}^M \left( \underbrace{A_{j, y^{j-1}}}_{\text{Scoring transitions}} + \underbrace{O_{y^j, x^j}}_{\text{Scoring observations}} \right)
 \end{aligned}$$

Maintain length-k prefix solutions	$\hat{Y}^k(T) = \left( \operatorname{argmax}_{y^{1:k-1}} F(y^{1:k-1} \oplus T, x^{1:k}) \right) \oplus T$
Recursively solve for length-(k+1) solutions	$  \begin{aligned}  \hat{Y}^{k+1}(T) &= \left( \operatorname{argmax}_{y^{1:k} \in \{\hat{Y}^k(T)\}_T} F(y^{1:k} \oplus T, x^{1:k+1}) \right) \oplus T \\  &= \left( \operatorname{argmax}_{y^{1:k} \in \{\hat{Y}^k(T)\}_T} F(y^{1:k}, x^{1:k}) + A_{T, y^k} + O_{T, x^{k+1}} \right) \oplus T  \end{aligned}  $
Predict via best length-M solution	$\operatorname{argmax}_y F(y, x) = \operatorname{argmax}_{y \in \{\hat{Y}^M(T)\}_T} F(y, x)$

**Solve:**  $\hat{Y}^2(V) = \left( \underset{y^1 \in \{\hat{Y}^1(T)\}_T}{\operatorname{argmax}} F(y^1, x^1) + A_{V, y^1} + O_{V, x^2} \right) \oplus V$


Store each  
 $\hat{Y}^1(T)$  &  $F(\hat{Y}^1(T), x)$



$\hat{Y}^1(T)$  is just  $T$

**Solve:**  $\hat{Y}^2(V) = \left( \underset{y^1 \in \{\hat{Y}^1(T)\}_T}{\operatorname{argmax}} F(y^1, x^1) + A_{V, y^1} + O_{V, x^2} \right) \oplus V$

Store each  
 $\hat{Y}^1(T)$  &  $F(\hat{Y}^1(T), x^1)$




$\hat{Y}^1(T)$  is just T

Ex:  $\hat{Y}^2(V) = (N, V)$

**Solve:**  $\hat{Y}^3(V) = \left( \operatorname{argmax}_{y^{1:2} \in \{\hat{Y}^2(T)\}_T} F(y^{1:2}, x^{1:2}) + A_{V, y^2} + O_{V, x^3} \right) \oplus V$

Store each  
 $\hat{Y}^1(T)$  &  $F(\hat{Y}^1(T), x^1)$

Store each  
 $\hat{Y}^2(Z)$  &  $F(\hat{Y}^2(Z), x)$



$\hat{Y}^1(Z)$  is just  $Z$


Ex:  $\hat{Y}^2(V) = (N, V)$

$$\textbf{Solve: } \hat{Y}^M(V) = \left( \operatorname{argmax}_{y^{1:M-1} \in \{\hat{Y}^M(T)\}_T} F(y^{1:M-1}, x^{1:M-1}) + A_{V, y^{M-1}} + O_{V, x^M} \right) \oplus V$$

Store each  
 $\hat{Y}^1(Z)$  &  $F(\hat{Y}^1(Z), x^1)$

Store each  
 $\hat{Y}^2(T)$  &  $F(\hat{Y}^2(T), x)$

Store each  
 $\hat{Y}^3(T)$  &  $F(\hat{Y}^3(T), x)$



$\hat{Y}^1(T)$  is just T

Ex:  $\hat{Y}^2(V) = (N, V)$

Ex:  $\hat{Y}^3(V) = (D, N, V)$



# Computing $P(y|x)$

- Viterbi doesn't compute  $P(y|x)$ 
  - Just maximizes the numerator  $F(y,x)$

$$P(y|x) = \frac{\exp\{F(y,x)\}}{\sum_{y'} \exp\{F(y',x)\}} \equiv \frac{1}{Z(x)} \exp\{F(y,x)\}$$

- Also need to compute  $Z(x)$ 
  - aka the “Partition Function”

$$Z(x) = \sum_{y'} \exp\{F(y',x)\}$$

# Computing Partition Function

- Naive approach is iterate over all  $y'$ 
  - Exponential time,  $L^M$  possible  $y'$ !

$$Z(x) = \sum_{y'} \exp\{F(y', x)\} \quad F(y, x) \equiv \sum_{j=1}^M \left( A_{y^j, y^{j-1}} + O_{y^j, x^j} \right)$$

- Notation:  $G^j(b, a) = \exp\{A_{b,a} + O_{b,x^j}\}$  Suppressing dependency on  $x$  for simpler notation


$$P(y | x) = \frac{1}{Z(x)} \prod_{j=1}^M G^j(y^j, y^{j-1})$$

$$Z(x) = \sum_{y'} \prod_{j=1}^M G^j(y'^j, y'^{j-1})$$

# Matrix Semiring

$$Z(x) = \sum_{y'} \prod_{j=1}^M G^j(y'^j, y'^{j-1})$$

$$G^j(b, a) = \exp \left\{ A_{b,a} + O_{a,x^j} \right\}$$



$$G^{1:2}(b, a) \equiv \sum_c G^2(b, c) G^1(c, a)$$




$$G^{i:j}(b, a) \equiv \text{[blue box } G^{i:j} \text{]} = \text{[blue box } G^j \text{]} \text{ [blue box } G^{j-1} \text{]} \dots \text{[blue box } G^{i+1} \text{]} \text{ [blue box } G^i \text{]}$$

# Path Counting Interpretation

- Interpretation  $G^1(b,a)$ 
  - L+1 start & end locations
  - Weight of path from 'a' to 'b' in step 1



- $G^{1:2}(b,a)$ 
  - Weight of all paths
    - Start in 'a' beginning of Step 1
    - End in 'b' after Step 2



# Computing Partition Function

- Consider Length-1 ( $M=1$ )

$$Z(x) = \sum_b G^1(b, \text{Start})$$

Sum column 'Start' of  $G^1$ !

- $M=2$

$$Z(x) = \sum_{a,b} G^2(b,a)G^1(a, \text{Start}) = \sum_b G^{1:2}(b, \text{Start})$$

Sum column 'Start' of  $G^{1:2}$ !

- General  $M$

Sum column 'Start' of  $G^{1:M}$ !

- Do  $M (L+1) \times (L+1)$  matrix computations to compute  $G^{1:M}$
- $Z(x) = \text{sum column 'Start' of } G^{1:M}$

$$G^{1:M} = G^M G^{M-1} \dots G^2 G^1$$

# Computing Partition Function

- Consider Length-1 ( $M=1$ )

$$Z(x) = \sum_l G^1(b, \text{Start})$$

start' of  $G^1$ !

- $M=2$

**Numerical Instability Issues!**  
(See Course Notes)

( $b, \text{Start}$ )

start' of  $G^{1:2}$ !

- General  $M$

sum column start' of  $G^{1:M}$ !

- Do  $M (L+1) \times (L+1)$  matrix computations to compute  $G^{1:M}$
- $Z(x) = \text{sum column 'Start' of } G^{1:M}$

$$G^{1:M} = G^M G^{M-1} \dots G^2 G^1$$

# Train via Gradient Descent

- Similar to Logistic Regression
  - Gradient Descent on negative log likelihood (log loss)

$$\operatorname{argmin}_{\Theta} \sum_{i=1}^N -\log P(y_i | x_i) = \operatorname{argmin}_{\Theta} \sum_{i=1}^N -F(y_i, x_i) + \log(Z(x_i))$$

$\Theta$  often used to denote all parameters of model

Harder to  
differentiate!

- First term is easy:

- Recall:

$$F(y, x) \equiv \sum_{j=1}^M \left( A_{y^j, y^{j-1}} + O_{y^j, x^j} \right)$$

$$\partial_{A_{ba}} - F(y, x) = - \sum_{j=1}^M 1_{[(y^j, y^{j-1}) = (b, a)]}$$

$$\partial_{O_{az}} - F(y, x) = - \sum_{j=1}^M 1_{[(y^j, x^j) = (a, z)]}$$

# Differentiating Log Partition

Lots of Chain Rule & Algebra!

$$\begin{aligned}
 \partial_{A_{ba}} \log(Z(x)) &= \frac{1}{Z(x)} \partial_{A_{ba}} Z(x) = \frac{1}{Z(x)} \partial_{A_{ba}} \sum_{y'} \exp\{F(y', x)\} \\
 &= \frac{1}{Z(x)} \sum_{y'} \partial_{A_{ba}} \exp\{F(y', x)\} \\
 &= \frac{1}{Z(x)} \sum_{y'} \exp\{F(y', x)\} \partial_{A_{ba}} F(y', x) = \sum_{y'} \frac{\exp\{F(y', x)\}}{Z(x)} \partial_{A_{ba}} F(y', x) \\
 &\stackrel{\text{Definition of } P(y' | x)}{=} \sum_{y'} P(y' | x) \partial_{A_{ba}} F(y', x) = \sum_{y'} \left[ P(y' | x) \sum_{j=1}^M 1_{[(y'^j, y'^{j-1}) = (b, a)]} \right] \\
 &= \sum_{j=1}^M \sum_{y'} P(y' | x) 1_{[(y'^j, y'^{j-1}) = (b, a)]} = \sum_{j=1}^M P(y^j = b, y^{j-1} = a | x)
 \end{aligned}$$

**Forward-Backward!**

↑  
Marginalize over all  $y'$



# Optimality Condition

$$\operatorname{argmin}_{\Theta} \sum_{i=1}^N -\log P(y_i | x_i) = \operatorname{argmin}_{\Theta} \sum_{i=1}^N -F(y_i, x_i) + \log(Z(x))$$

- Consider one parameter:

$$\partial_{A_{ba}} \sum_{i=1}^N -F(y_i, x_i) = -\sum_{i=1}^N \sum_{j=1}^{M_i} 1_{[(y_i^j, y_i^{j-1})=(b,a)]} \quad \partial_{A_{ba}} \sum_{i=1}^N \log(Z(x)) = \sum_{i=1}^N \sum_{j=1}^{M_i} P(y_i^j = b, y_i^{j-1} = a | x_i)$$

- Optimality condition:

$$\sum_{i=1}^N \sum_{j=1}^{M_i} 1_{[(y_i^j, y_i^{j-1})=(b,a)]} = \sum_{i=1}^N \sum_{j=1}^{M_i} P(y_i^j = b, y_i^{j-1} = a | x_i)$$

- **Frequency counts = Cond. expectation on training data!**
  - Holds for each component of the model
  - Each component is a “log-linear” model and requires gradient desc.

# Forward-Backward for CRFs

$$\alpha^1(a) = G^1(a, \text{Start})$$

$$\alpha^j(a) = \sum_{a'} \alpha^{j-1}(a') G^j(a, a')$$

$$\beta^M(b) = 1$$


$$\beta^j(b) = \sum_{b'} \beta^{j+1}(b') G^{j+1}(b', b)$$

$$P(y^j = b, y^{j-1} = a \mid x) = \frac{\alpha^{j-1}(a) G^j(b, a) \beta^j(b)}{Z(x)}$$

$$Z(x) = \sum_{y'} \exp\{F(y', x)\} \quad F(y, x) \equiv \sum_{j=1}^M (A_{y^j, y^{j-1}} + O_{y^j, x^j}) \quad G^j(b, a) = \exp\{A_{b, a} + O_{b, x^j}\}$$

# Path Interpretation

Total Weight of paths from "Start" to "V" in 3<sup>rd</sup> step



$\beta$  just does it backwards

# Matrix Formulation

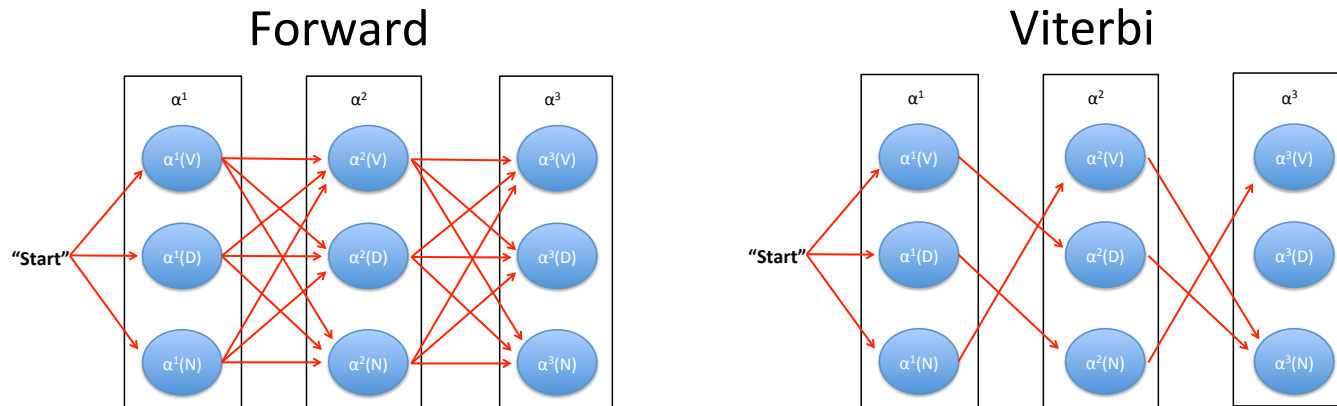
- Use Matrices!
- Fast to compute!
- Easy to implement!

$$\alpha^2 = G^2 \alpha^1$$

$$\beta^6 = (G^2)^T \beta^5$$

# Path Interpretation:

## Forward-Backward vs Viterbi



- Forward (and Backward) sums over all paths
  - Computes expectation of reaching each state
  - E.g., total (un-normalized) probability of  $y^3=\text{Verb}$  over all possible  $y^{1:2}$
- Viterbi only keeps the best path
  - Computes best possible path to reaching each state
  - E.g., single highest probability setting of  $y^{1:3}$  such that  $y^3=\text{Verb}$

# Summary: Training CRFs

- Similar optimality condition as HMMs:
    - Match frequency counts of model components!
- $$\sum_{i=1}^N \sum_{j=1}^{M_i} 1_{[(y_i^j, y_i^{j-1})=(b,a)]} = \sum_{i=1}^N \sum_{j=1}^{M_i} P(y_i^j = b, y_i^{j-1} = a | x_i)$$
- Except HMMs can just set the model using counts
  - CRFs need to do gradient descent to match counts
- Run Forward-Backward for expectation
    - Just like HMMs as well

# Summary: CRFs

- Log-Linear Sequential Model:

$$P(y | x) = \frac{\exp\{F(y, x)\}}{Z(x)}$$
$$F(y, x) \equiv \sum_{j=1}^M (A_{y^j, y^{j-1}} + O_{y^j, x^j})$$
$$Z(x) = \sum_{y'} \exp\{F(y', x)\}$$

- Same #parameters as HMMs
  - Focused on learning  $P(y|x)$
  - Prediction via Viterbi
  - Gradient Descent via Forward-Backward

# Next Lecture

- More General Formulation of CRFs
  - More concise notation
    - Matches logistic regression notation
    - Matches course notes (later this week)
  - Easier to reason about for implementation
- General Structured Prediction